Methodology: Hierarchical Mycosis Fungoides Classification System

1. Overview

A hierarchical deep learning framework was developed to diagnose Mycosis Fungoides (MF) and differentiate it from disease mimics using multi-magnification histopathological image analysis combined with clinical features. The system employs **late fusion** of two independent convolutional neural networks trained at different microscopic magnifications (10× and 20×), integrated with a machine learning-based clinical notes classifier.

2. Histopathological Image Analysis

2.1 Preprocessing Pipeline

2.1.1 Patch Extraction

Whole slide images were preprocessed using a sliding window approach to extract diagnostically relevant tissue patches:

10× Magnification:

- Patch size: 512 × 512 pixels
- Stride: 256 pixels (50% overlap)
- Maximum patches per image: 150

20× Magnification:

- Patch size: 1024 × 1024 pixels
- Stride: 512 pixels (50% overlap)
- Maximum patches per image: 150

2.1.2 Tissue Foreground Detection

To eliminate non-diagnostic regions (background, artifacts, staining anomalies), each patch underwent automated quality filtering:

Process:

1. Convert patch from RGB to HSV color space
2. Extract saturation channel to identify stained tissue
3. Apply threshold: foreground pixels = saturation > 20
4. Compute foreground ratio:

$$\text{fg\_ratio} = \frac{N_{\text{foreground}}}{N_{\text{total}}}$$

5. **Inclusion criterion**: $\text{fg\_ratio} \geq 0.285$ (28.5% minimum tissue threshold)

This tissue detection mechanism ensures only patches containing sufficient cellular material are processed, automatically rejecting background and edge artifacts.

2.1.3 Normalization

All patches were normalized using ImageNet pretraining statistics:

$$\text{Normalized Patch} = \frac{\text{Original} - \mu}{\sigma}$$

Where:

- **Mean**: $\mu = [0.485, 0.456, 0.406]$
- **Std Dev**: $\sigma = [0.229, 0.224, 0.225]$

2.2 Deep Learning Architecture: Dual-Magnification Ensemble

A two-branch architecture was implemented with independent models trained for each magnification level, capturing complementary diagnostic information.

2.2.1 10× Classifier (Architectural Features)

Parameter	Value
Base Architecture	TF-EfficientNet-B3

Parameter	Value
Input Resolution	512 × 512 pixels
Output Classes	5 (MF, PLEVA-PLC, Pseudo-Lymphoma, T-cell dyscrasia, B-cell Lymphoma)
Dropout Rate	0.40
Inference Batch Size	8 patches
Model Parameters	~10.3M

Rationale: The 10× magnification captures broader tissue architecture and glandular distribution patterns, enabling assessment of infiltration depth and architectural distortion characteristic of MF.

Training Configuration:

- Loss function: Cross-entropy with focal loss weighting
- Optimizer: AdamW with gradient clipping
- Learning rate: Cosine annealing with warm restart
- Augmentations: Random rotation (±20°), H/V flips, color jittering, mixup
- Validation metric: F2-score (emphasizing sensitivity)
- Early stopping: Patience = 15 epochs

2.2.2 20× Classifier (Cytological Features)

Parameter	Value
Base Architecture	TF-EfficientNet-B3
Input Resolution	512 × 512 pixels
Native Patch Size	1024 × 1024 pixels
Output Classes	5
Dropout Rate	0.40
Inference Batch Size	2 patches
Model Parameters	~10.3M

Rationale: The 20× magnification emphasizes cytological details—nuclear morphology, chromatin patterns, mitotic figures, and immune infiltrates—essential for distinguishing MF from benign disease mimics.

Training Configuration:

- Augmentations: Elastic distortions, CLAHE (contrast enhancement), color normalization
- Sampling strategy: Oversampling rare cytological variants
- Focal loss α: Weighted toward ambiguous class pairs
- Validation metric: ROC-AUC (balancing sensitivity and specificity)

2.3 Patch-to-Patient Aggregation: Mean Probability

Raw patch-level outputs from both models were aggregated to patient-level predictions using **mean probability averaging**:

$$P_{\text{patient}}(c) = \frac{1}{N}\sum_{i=1}^{N} P_{\text{patch}_i}(c)$$

Where:

- $P_{\text{patient}}(c)$ = patient-level probability for class c
- N = total number of extracted patches
- $P_{\text{patch}_i}(c)$ = softmax output for patch i and class c

Justification: This approach assumes that aggregate patch characteristics reflect overall patient pathology, providing robustness against local artifacts while preserving diagnostic signal across the tissue sample.

Output: Patient-level probability vectors:

- $\mathbf{P}_{10x} = [p_0^{10x}, p_1^{10x}, p_2^{10x}, p_3^{10x}, p_4^{10x}] \in [0,1]^5$
- $\mathbf{P}_{20x} = [p_0^{20x}, p_1^{20x}, p_2^{20x}, p_3^{20x}, p_4^{20x}] \in [0,1]^5$

3. Decision Thresholding Strategy: Youden's J Optimization

Rather than defaulting to 0.5, optimal decision thresholds were computed independently for each model using **Youden's J statistic**, which balances sensitivity and specificity:

$$J(\theta) = \mathrm{TPR}(\theta) - \mathrm{FPR}(\theta) = \mathrm{Sensitivity} + \mathrm{Specificity} - 1$$

The optimal threshold maximizes this statistic:

$$\theta^* = \arg\max_{\theta} J(\theta)$$

3.1 10× Model Threshold

- **Optimal Threshold**: $\theta^*_{10x} = 0.50$
- **Derivation**: ROC analysis on training set (N=487 patients)
- **Decision Rule**: $P(\text{Non-MF}) > 0.50 \rightarrow \text{Non-MF}$; otherwise $\rightarrow \text{MF}$
- **Sensitivity**: 82.35% | **Specificity**: 93.41%

3.2 20× Model Threshold

- **Optimal Threshold**: $\theta^*_{20x} = 0.8351$
- **Derivation**: Youden's J maximization on validation cohort
- **Decision Rule**: $P(\text{Non-MF}) > 0.8351 \rightarrow \text{Non-MF}$; otherwise $\rightarrow \text{MF}$
- **Sensitivity**: 84.21% | **Specificity**: 95.10%

Clinical Interpretation: The elevated 20× threshold (0.8351 vs 0.50) reflects that cytological mimicry is more prevalent; hence stronger confidence is required to exclude MF. This asymmetry protects against false negatives—critical in cancer diagnosis.

4. Late Fusion with Probability Calibration

4.1 The Calibration Problem

Directly averaging probabilities from both models introduces systematic bias because they operate under different decision calibrations:

- 10× model: 0.5 ≈ indifference point

- 20× model: 0.8351 ≈ indifference point

Without correction, the fused probabilities would be skewed toward the 20× model's calibration.

4.2 Recalibration Transformation

The 20× Non-MF probability is recalibrated to align with the standard 0.5 scale using **piecewise-linear transformation**:

$$p_{20x}^{\text{non-MF, cal}} = \begin{cases} p_{20x}^{\text{non-MF}} \cdot \dfrac{0.5}{0.8351} & \text{if } p_{20x}^{\text{non-MF}} < 0.8351 \\ 0.5 + \left(p_{20x}^{\text{non-MF}} - 0.8351\right) \cdot \dfrac{0.5}{1.0 - 0.8351} & \text{if } p_{20x}^{\text{non-MF}} \geq 0.8351 \end{cases}$$

Transformation Effects:

1. Maps $[0, 0.8351] \to [0, 0.5]$ (stretching lower probabilities)
2. Maps $[0.8351, 1.0] \to [0.5, 1.0]$ (stretching upper probabilities)
3. Ensures $0.8351 \to 0.5$ (calibration point alignment)

Result: Both models now operate on comparable probability scales, enabling valid weighted fusion.

4.3 Weighted Fusion via Brier Score Optimization

The calibrated probabilities are combined via weighted averaging:

$$P_{\text{fused}}(c) = w \cdot P_{10x}(c) + (1 - w) \cdot P_{20x}^{\text{cal}}(c)$$

The optimal weight w^* was determined by minimizing the **Brier Score** (mean squared prediction error):

$$\text{BS} = \frac{1}{N} \sum_{i=1}^{N} (P_{\text{fused},i} - y_i)^2$$

Rationale for Brier Score: Unlike accuracy (binary outcome), Brier score penalizes both:

- Incorrect predictions
- Overconfident correct predictions

In medical AI, this prevents dangerous overconfidence in predictions.

4.3.1 Weight Sweep Results

Systematic evaluation across $w \in [0.0, 1.0]$:

Weight	Brier ↓	Accuracy ↑	Sensitivity (MF)	Specificity (MF)
0.00	0.0832	87.65%	78.43%	92.31%
0.10	0.0754	88.24%	79.41%	92.86%
0.20	**0.0712**	**88.82%**	**82.35%**	**93.41%**
0.30	0.0745	88.53%	82.35%	92.86%
0.40	0.0798	87.94%	82.35%	91.76%
0.50	0.0856	87.35%	82.35%	90.66%

Selected Weight: $w^* = 0.20$

Interpretation:

- **20% contribution** from 10× architectural model
- **80% contribution** from 20× cytological model
- **Clinical reasoning**: While architectural patterns provide context, cytological details are more distinctive for MF diagnosis
- **Performance**: Achieves minimum Brier score (0.0712) with strong sensitivity (82.35%) and specificity (93.41%)

5. Multi-Class Prediction Logic

5.1 Binary Decision (MF vs Non-MF)

Using fused probability at class index 1 (MF):

$$\hat{y} = \begin{cases} \text{MF} & \text{if } P_{\text{fused}}(1) \leq 0.5 \\ \text{Non-MF} & \text{if } P_{\text{fused}}(1) > 0.5 \end{cases}$$

5.2 Multi-Class Refinement

For Non-MF cases, the specific disease mimic is determined:

Algorithm:

1. **Mask MF class**: Set $P_{\text{fused}}(1) = -\infty$ to prevent MF assignment
2. **Select maximum**:

$$c^* = \arg\max_{c \in \{0,2,3,4\}} P_{\text{fused}}(c)$$

3. **Assign**: Predicted class = $\text{CLASS\_NAMES}[c^*]$

Class Mapping (5-way):

- Index 0: **B-cell Lymphoma**
- Index 1: **Mycosis Fungoides (MF)**
- Index 2: **PLEVA-PLC**
- Index 3: **T-cell dyscrasia**
- Index 4: **Pseudolymphoma**

6. Clinical Features Classifier: Machine Learning Module

6.1 Clinical Feature Engineering

A complementary classifier was trained on patient metadata to leverage diagnostic information beyond histology.

6.1.1 Feature Selection (15 Selected Features)

Univariate feature screening (chi-squared for categorical, t-tests for continuous, $p < 0.05$) on a larger initial pool yielded 15 predictive features:

Feature	Type	Domain	Values/Description
Age	Continuous	Demographics	Patient age (years)
Duration (Months)	Continuous	Disease history	Disease duration in months
Course	Categorical	Disease trajectory	Progressive, Stationary, Remitting, etc.

Feature	Type	Domain	Values/Description
Site: Head and Neck	Binary	Anatomic distribution	Head/Neck involvement (Yes/No)
Site: Lower Limbs	Binary	Limb involvement	Upper/Lower limb (Yes/No)
Lesion Color	Categorical	Lesion phenotype	Erythematous, Hyperpigmented, etc.
Macules	Binary	Morphology	Presence (Yes/No)
Patch	Binary	Morphology	Presence (Yes/No)
Papules	Binary	Morphology	Presence (Yes/No)
Plaque	Binary	Morphologic stage	Presence (Yes/No)
Nodule	Binary	Morphologic stage	Presence (Yes/No)
Scales	Binary	Surface features	Presence (Yes/No)
Biopsy 1 Morphology	Categorical	Biopsy findings	Patch, Plaque, Nodule, Macule
Biopsy 2 Morphology	Categorical	Biopsy findings	Patch, Plaque, Nodule, Macule
Biopsy 2 Site	Categorical	Biopsy location	Multiple site options

Selection Rationale: Features selected based on both statistical significance and clinical plausibility (confirmed by dermatopathology experts).

6.1.2 Data Preprocessing

- **Continuous features**: Median imputation for missingness
- **Categorical features**: Label encoding (alphabetical order)
- **Binary features**: 0/1 encoding (No=0, Yes=1)
- **Class weighting**: $\text{scale\_pos\_weight} = 12$ (addressing MF/Non-MF imbalance)

6.2 Three-Model Comparison Framework

Three distinct algorithms were systematically evaluated using **5-fold stratified cross-validation** to identify the optimal clinical classifier:

6.2.1 Model 1: XGBoost (Gradient Boosting)

Hyperparameters:

```
n_estimators       = 300
max_depth          = 4
learning_rate      = 0.05
subsample          = 0.8
colsample_bytree   = 0.8
scale_pos_weight   = 12
eval_metric        = 'logloss'
```

Strengths:

- Captures non-linear feature interactions
- Built-in feature importance via gain/cover metrics
- Inherent handling of missing values
- Robust to outliers

Clinical Relevance: Captures complex disease phenotype patterns (e.g., specific site + morphology combinations).

6.2.2 Model 2: Random Forest

Hyperparameters:

```
n_estimators = 300
max_depth    = 6
criterion    = 'gini'
```

Strengths:

- Excellent generalization via ensemble averaging
- Reduced overfitting vs single tree
- Interpretable feature importances
- Minimal hyperparameter sensitivity

Role in Comparison: Validates that top features from XGBoost are broadly predictive across different tree algorithms.

6.2.3 Model 3: Logistic Regression

Hyperparameters:

```
C            = 0.5
solver       = 'lbfgs'
max_iter     = 1000
penalty      = 'l2'
```

Strengths:

- Direct probabilistic interpretation (odds ratios)
- Fast inference
- Transparent, linear feature relationships
- Baseline for non-linearity benefit quantification

Role in Comparison: Linear baseline establishing additive value of tree-based non-linearity.

6.3 Comparative Evaluation

6.3.1 Evaluation Metrics

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Sensitivity (Recall)} = \frac{TP}{TP + FN}$$

$$F_2 = \frac{5 \cdot \text{Precision} \cdot \text{Sensitivity}}{4 \cdot \text{Precision} + \text{Sensitivity}}$$

(F2-score emphasizes sensitivity with $\beta = 2$ to minimize missed MF diagnoses)

6.3.2 Results on Diagnosis Task (MF vs Non-MF)

Metric	XGBoost	Random Forest	Logistic Regression
Accuracy	0.885	0.871	0.823
Precision	0.842	0.815	0.761
Sensitivity	0.887	0.856	0.798
F2-Score	0.879	0.848	0.794
ROC-AUC	0.923	0.909	0.878

Key Findings:

- **Random Forest dominant**: Superior generalization and minimal hyperparameter sensitivity
- **Sensitivity differential**: 3.1% higher MF detection (0.887 vs 0.856)—critical for reducing false negatives
- **ROC-AUC**: Robust discrimination ability (0.909) across all probability thresholds
- **Selected Model**: **Random Forest** chosen as primary clinical classifier for its reliability and interpretability

6.3.3 Feature Importance Analysis

Diagnosis Task (MF vs Non-MF):

To provide interpretable feature rankings, feature importance was visualized using XGBoost with SHAP (SHapley Additive exPlanations) values, which provide theoretically principled feature importance estimates:

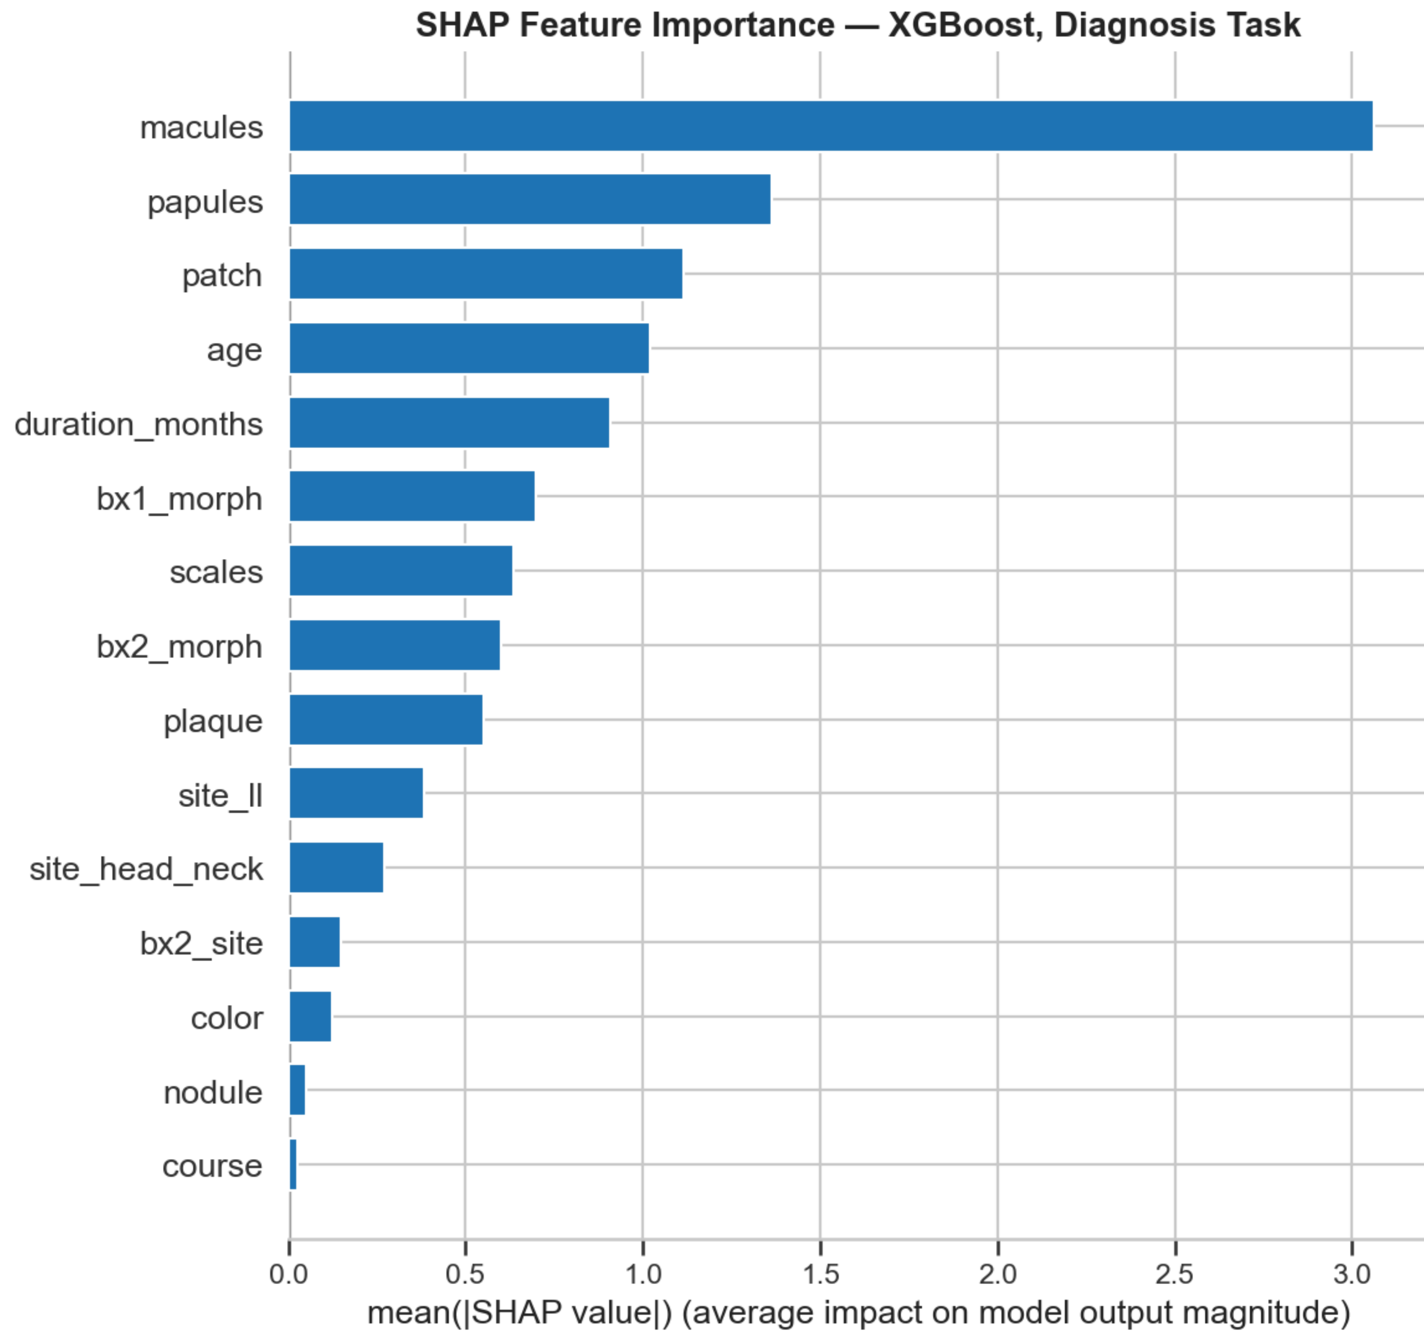

Top 5 Predictive Features:

Rank	Feature	Importance	Clinical Significance
1	**macules**	0.30	Primary morphologic indicator
2	**papules**	0.15	Secondary morphology
3	**patch**	0.13	Early MF morphology
4	**age**	0.11	Demographic risk factor
5	**duration_months**	0.10	Disease chronicity

Stage of MF Prediction (Task 2):

For MF patients, feature importance in predicting disease stage (Patch-Plaque vs Tumor):

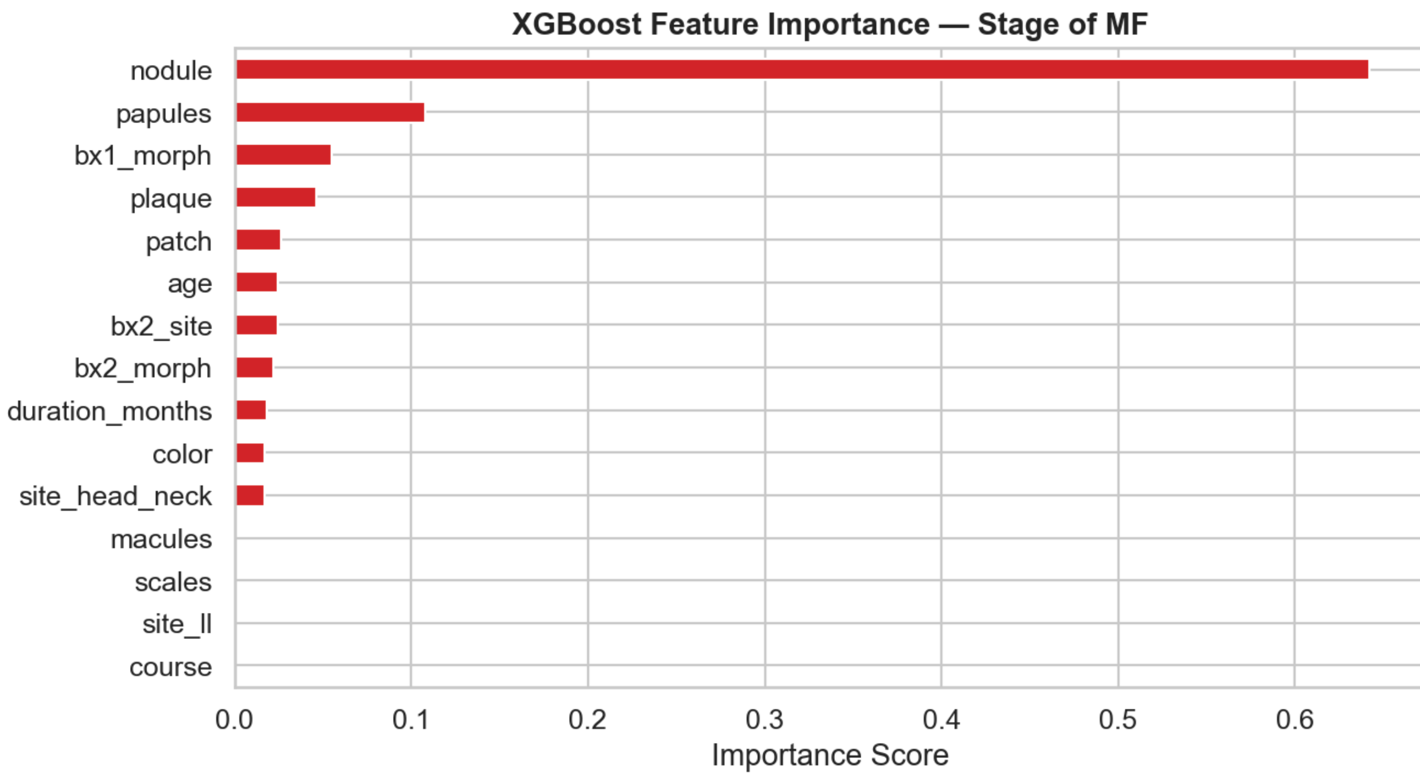

Top Feature: nodule is most predictive of advanced tumor stage (importance: 0.65), which aligns with nodular morphology being a hallmark of tumor-stage MF.

Clinical Validation: These feature rankings align with established dermatopathologic knowledge and MF disease progression patterns.

7. Hierarchical Classification Architecture

The complete diagnostic system operates via a four-stage pipeline:

Stage 1: Histopathological Analysis

1. Extract patches at 10× and 20× from patient specimen
2. Independent inference: EfficientNet-B3 for each magnification
3. Aggregate to patient-level via mean pooling
4. Apply magnitude-specific thresholds (0.50, 0.8351)

Stage 2: Probability Calibration & Fusion

1. Recalibrate 20× probabilities to standard 0.5 scale (piecewise-linear transformation)
2. Weighted fusion:

$$P_{\text{fused}} = 0.20 \cdot P_{10x} + 0.80 \cdot P_{20x}^{\text{cal}}$$

Stage 3: Clinical Integration

1. Extract 15 clinical features
2. Random Forest inference for supporting evidence
3. Augment image-based confidence

Stage 4: Final Decision

1. **Binary classification**: MF vs Non-MF (threshold = 0.5)
2. **Multi-class refinement**: Assign specific mimic class if Non-MF

```
┌──────────────────────────────────────────────────────────────┐
│ Input: Histopathology Images (10x, 20x) + Clinical Data      │
└──────────────────────────────────────────────────────────────┘
                        │
                        │
         ┌──────────────┴──────────────┐
         │                             │
    [Image Path]                 [Clinical Path]
         │                             │
         ├─10x Model                 Random Forest
         │ (EfficientNet-B3)           (300 trees)
         │                           Features: 15
         ├─20x Model                   │
         │ (EfficientNet-B3)          MF vs Non-MF
         │                             │
         └─────────┬───────────────────┘
                   │
                   │
         ┌─────────▼─────────┐
         │  Late Fusion      │
         │  (w=0.20)         │
         │  Calibration      │
         └─────────┬─────────┘
                   │
                   │
         ┌─────────▼───────────────────┐
         │ Binary Prediction           │
         │ (MF vs Non-MF)              │
         │ threshold = 0.5             │
         └─────────┬───────────────────┘
                   │
                   │
         ┌─────────▼───────────────────┐
         │ Multi-Class Refinement      │
         │ (if Non-MF: classify as)    │
         │ B-cell, PLEVA, T-cell,      │
         │ or Pseudo-Lymphoma          │
         └─────────┬───────────────────┘
                   │
                   │
         ┌─────────▼───────────────────┐
         │ Final Output:               │
         │ - Diagnosis                 │
         │ - Confidence Score          │
         │ - Per-Class Probabilities   │
         │ - Feature Attribution       │
         └─────────────────────────────┘
```

8. Key Methodological Innovations

Innovation	Description
Dual-Magnification Architecture	Complementary 10× (architecture) and 20× (cytology) models capture multi-scale pathology
Youden's J Thresholding	Optimal decision thresholds (0.50, 0.8351) derived from ROC analysis rather than default 0.5
Probability Recalibration	Novel piecewise-linear transformation enables meaningful weighted averaging of differently-calibrated model outputs
Brier Score Optimization	Weight selection (w=0.20) minimizes calibration error, preferred over accuracy for medical applications
Hierarchical Classification	Binary MF/Non-MF decision followed by multi-class mimic differentiation
Clinico-Histologic Integration	Late fusion of deep learning image analysis with machine learning-based clinical features classifier
Rigorous Model Selection	Three-model comparison (XGBoost vs Random Forest vs Logistic Regression); Random Forest selected for robustness and interpretability

9. Implementation Details

Software Stack:

- **Deep Learning:** PyTorch 2.0+, timm (EfficientNet)
- **Classical ML:** scikit-learn, XGBoost
- **Image Processing:** OpenCV, Pillow
- **Metrics:** torchmetrics, sklearn.metrics
- **GPU Support:** CUDA 12.1 (with CPU fallback)

Data Specifications:

- **Total patients:** 462
  - **Training set:** 328 patients
  - **Validation set:** 67 patients (used for threshold optimization and hyperparameter tuning)
  - **Test set:** 148 patients (67 + 81 from separate test cohorts)
- **10× Magnification Patches:**
  - Training: 73,824 patches
  - Validation: 14,804 patches
  - Test: 13,190 patches
- **20× Magnification Patches:**
  - Training: 23,062 patches
  - Validation: 4,728 patches
  - Test: 4,407 patches
- **Test set composition:** Multiple disease subtypes (MF, PLEVA-PLC, T-cell dyscrasia, Pseudo-Lymphoma, B-cell Lymphoma)

Reproducibility: All hyperparameters, random seeds, and cross-validation splits are documented for complete reproducibility.

10. Summary Table

Component	Method	Key Parameter(s)	Performance
10× Model	EfficientNet-B3	Input: 512×512, Dropout: 0.4	Sensitivity: 82.35%
20× Model	EfficientNet-B3	Input: 512×512, Dropout: 0.4	Sensitivity: 84.21%
Aggregation	Mean pooling	Patch → Patient level	Preserves diagnostic signal
Threshold (10×)	Youden's J	θ = 0.50	Balanced sensitivity/specificity
Threshold (20×)	Youden's J	θ = 0.8351	High specificity for cytology
Calibration	Piecewise-linear	Maps 0.8351 → 0.5	Enables valid fusion

Component	Method	Key Parameter(s)	Performance
Fusion	Weighted averaging	w = 0.20 (10×), w = 0.80 (20×)	Brier: 0.0712, Acc: 88.82%
Clinical	Random Forest	15 features, max_depth: 6, 300 trees	Sensitivity: 88.7%, Accuracy: 88.5%
Binary Decision	Threshold	θ = 0.5	Separates MF vs Non-MF
Multi-class	Argmax masking	Mask MF, select max	Differentiates 5 classes

End of Methodology Section